

Paper SM11

Code that writes code that writes code?!

Creating ADaMs in R with {admiralroche}

Gordon Miller, Roche Products Ltd, Welwyn Garden City, UK
Edoardo Mancini, Roche Products Ltd, Welwyn Garden City, UK

PHUSE EU Connect 2026, Glasgow, UK

1 Abstract

Writing template code is just as much an art as it is a skill, requiring balancing solid starting points with flexibility for custom adaptations. For developers, maintaining templates while ensuring consistency and eliminating repetitive tasks is a constant struggle.

Over the past years, Roche’s fast transition to R for ADaM programming challenged the {admiralroche} team to rethink our process, requiring constant evolution to provide fit for purpose templates. In this paper, we walk readers through our journey of levelling up our templates, starting by exploring how and why we initially set up the {admiralroche} package extension and why we then centralised template maintenance by deconstructing templates into modular blocks (code that writes code). Finally, we unveil our latest frontier: dynamic ADaM template generation. By passing ADaM specifications to our package, we automatically create and assemble the right blocks into a customized template script (code that writes code that writes code!). We will discuss our challenges and learnings, sharing metrics on reduced maintenance and time saved. Join us to see why, in the quest for the ultimate ADaM template, deterministic automation trumps probabilistic AI.

2 Introduction

The {admiral} package and its wider “pharmaverse” ecosystem [pharmaverse.org] are the **de facto** standard tools to create Analysis Data Model (ADaM) datasets in R. {admiral} deliberately does not attempt to be a complete, drop-in solution; rather, it provides composable, reusable derivation functions and leaves the programmer to decide *which* functions to use, *how* to use them, *in what order* to do so, and most importantly, *where* to intersperse sponsor/study-specific code. This is by design, since the fine details of how any given company or study implements the ADaM standards inevitably vary.

This paradigm has led some sponsors to build company-specific extensions on top of {admiral}, which can both:

- Provide custom functions to implement company-specific derivation choices;
- Make available a library of ADaM template scripts that implement those choices in a consistent, repeatable way.

For a presentation and paper discussing Roche’s own journey setting up `{admiralroche}`, the interested reader can consult Ross Farrugia’s “OS05: `{admiral}` – From Open Source Through to Company Implementation” presentation and paper from the PHUSE US Connect 2024 conference. Building on this material, this paper and the talk it accompanies tell the story of how the way we provide ADaM template scripts through `{admiralroche}` has evolved over time, moving through three distinct stages that give the talk its title:

1. **Code.** A library of static, human-written ADaM template scripts, one per domain, distributed inside `{admiralroche}` itself.
2. **Code that writes code.** An internal R package, `{scriptr}`, that abstracts the common building blocks shared across those templates into reusable, parameterised “chunks” which are then assembled programmatically rather than copy-pasted.
3. **Code that writes code that writes code.** A companion package, `{admiralrochetools}`, that combines `{scriptr}`’s assembly engine with a short interactive questionnaire and a study’s own specification, so that running a single function produces a template already tailored to that study’s derivation choices, with no manual search-and-replace required.

3 Background: Why `{admiralroche}` and ADaM templates?

3.1 `{admiralroche}` covers the ~10% of Roche’s ADaM standards that `{admiral}` does not

Heuristically, it has been observed that for most given ADaM datasets at Roche (and most likely for other sponsors too), around 80% of the variable and parameter derivations are covered by functions from the `{admiral}` package, such as `derive_vars_dt()`, `derive_vars_merged()`, `derive_param_exist_flag()`. A further ~10% is manually programmed code that is genuinely specific to an individual study, and has no business being standardised into any shared package. The remaining ~10% is where `{admiralroche}` comes in: it reflects Roche’s own interpretation of the ADaM standards, for example, which imputation techniques to apply, sponsor-specific variables and parameters, helper code to interact with Roche’s chosen statistical computing environment (SCE), and glue code for other pharmaverse packages such as `{metatools}`, `{metacore}` and `{xportr}` which Roche has chosen to use as part of its ADaM pipeline.




Figure 1: Code split across ADaM programs at Roche

3.2 Accessing ADaM templates

Even before any sponsor extension enters the picture, `{admiral}` already makes starting a new ADaM program straightforward, via `use_ad_template()`. Calling it with a domain name and `package = "admiral"` generates a ready-made, generic template script for that domain directly into the user’s project:

```
library(admiral)

admiral::use_ad_template(
  "adsl",
```

```
package = "admiral"
)
```

However, these templates are simply generic starting points. As mentioned in the previous section, most sponsors will likely wish to provide more customised templates which are closer to their implementation of ADaM standards, work better within their SCE, etc. If a sponsor has created a sponsor-specific extension package like {admiralroche}, they can distribute their own versions of ADaM templates by leveraging the same `use_ad_template()` function and updating the `package` argument.

```
library(admiral)
library(admiralroche)

admiral::use_ad_template(
  "adsl",
  package = "admiralroche"
)
```

3.3 Features of sponsor ADaM templates

As expected, {admiralroche}'s ADaM templates showcase a healthy mix of standard {admiral} code customised to fit Roche's requirements, calls to {admiralroche} functions, and placeholders for study-specific code. Below are examples of each from the ADaM template.

Helper code to read in data cut date:

```
## Datacut date ----
# Get cut-off date from ATLAS activity.json meta file
activity_json <- read_json("metadain/atlas/activity.json")
cut_date <- activity_json$activity_ccod
# pass in datacut date (in form as.Date("2020-12-31") for example)
adsl_dcut <- dm %>%
  # datacut date (DCUTDT)
  mutate(DCUTDT = as.Date(cut_date))
```

Leveraging an {admiral} function but using Roche-specific arguments to derive treatment start date:

```
# treatment start date (TRTSDTM)
derive_vars_merged(
  dataset_add = ex_ext,
  filter_add = (EXDOSE > 0 | (EXDOSE == 0 & str_detect(EXTRT, "PLACEBO"))) &
    !is.na(EXSTDTM),
  new_vars = exprs(TRTSDTM = EXSTDTM),
  order = exprs(EXSTDTM, EXSEQ),
  mode = "first",
  by_vars = get_admiral_option("subject_keys")
) %>%
```

Using {admiralroche}-specific functions to derive Roche ADaM variables:

```
# age group 1 (AGEGR1)
admiralroche::derive_var_agegr1() %>%
# age group 2 (AGEGR2)
```

```
admiralroche::derive_var_agegr2() %>%
```

Study-specific treatment arm derivation placeholders:

```
# treatment variables (TRT01P/TRT01A)
mutate(
  TRT01P = ARM,
  TRT01A = ACTARM
) %>%
```

4 Stage 1: *Code* - static ADaM templates

Throughout its first releases, {admiralroche} put the `use_ad_template()` mechanism introduced in Section 3 to direct use: it shipped a static, human-written ADaM template script for every domain covered by Roche’s global standards - ADSL, ADAE, ADVS, ADLB, and so on — each one hard-coding Roche’s standard interpretation of that domain’s derivations and mixing together the {admiral} calls, {admiralroche} helper functions, and study-specific placeholders illustrated in Section 3.3. A programmer starting a new study would call the appropriate `use_*_template()` function, receive a copy of the relevant script in their project, and edit it by hand to fit the study at hand.

4.1 Pros and cons of the static template approach

This approach was the simplest possible way to package company-specific ADaM logic, and its strengths followed directly from that simplicity:

- **Simple to maintain for developers.** A template was just a plain R script; adding or updating one required no tooling beyond an editor and {admiralroche}’s usual code review process.
- **Simple for users to access and use.** A single call to `use_ad_template("<adam_name>", package = "admiralroche")` was all a programmer ever needed to retrieve a Roche-specific starting point.
- **The easiest way to initially set up {admiralroche}.** Static templates let the package reach its first users quickly, without first having to design any abstraction layer.

The cost of that simplicity, however, grew with the size of the template library:

- **Difficult to cater for many different options, especially if complex.** Each template hard-coded one interpretation of a derivation - usually the most popular one from Roche’s global standards. User who needed a different option had to find and edit the relevant lines by hand, with limited guidance in the form of a comment to help guide the way.
- **Duplication of code for sections common to all templates.** By the time {admiralroche} had accumulated more than twenty ADaM templates, several recurring patterns, for example, the logic for deriving AENDTM/AENDT consistently across the ADCM and ADMH templates, or standard subject-level flag derivations repeated across nearly every template, existed as near-identical, hand-copied blocks of code in many files at once.

This second point turned out to be the more serious limitation in practice. Any refactor, bug fix, or change in a Roche standard had to be propagated by hand into every template that contained the affected block, and so there was always a risk that a fix may land in some templates but be missed in others.




Figure 2: Duplication risks and comments pointing users towards potential updates.

5 Stage 2: *Code that writes code* - Abstracting away the ADaM templates

5.1 The problem: duplication without a shared abstraction

As hinted, the static-template approach was a good starting point, but had no mechanism for saying “this block of code is used, with minor variations, in these six templates - treat it as one thing.” Every template was an independent, flat R script, so the *only* way to reuse logic across templates was to copy it. To circumvent this issue, the {admiralroche} team developed an internal helper package, {scriptr}. This package introduced a proper abstraction layer between “a reusable unit of ADaM logic” and “a finished script” so that changes to the underlying logic only ever needed to be made in one place.

5.2 {scriptr}’s core data model: “chunks” and “calls”

{scriptr} organises code around two first-class objects, referred to internally as **chunks** and **calls**. A chunk is a general-purpose, reusable block of R code (i.e. anything from a single derivation call to a short multi-line sequence) registered once via `add_chunk()` and retrieved later via `get_chunk()`. A call is the more specialised case of a chunk that wraps exactly one function invocation (typically an {admiral} derivation function), registered via `add_call()` and retrieved via `get_call()`. Both chunks and calls can also be registered as *templates* — via `add_chunk_template()` and `add_call_template()` respectively — meaning they contain placeholders to be filled in at assembly time rather than fixed, literal code.

```

library(scriptr)
library(admiral)

add_chunk(
  id = "load_ae",

```

```

description = "Read the AE SDTM domain",
{
  ae <- read_parquet("data/sdtm/ae.parquet")
}
)

add_call(
  id = "derive_aendt",
  description = "Derive the AE end date (AENDT) for a single template",
  derive_vars_dt(
    new_vars_prefix = "AEN",
    dtc = AEENDTC,
    highest_imputation = "M",
    date_imputation = "first"
  )
)

```

As written, `derive_aendt` is only useful for the one template that happens to need `AENDTM/AENDT` derived with “first”-day imputation - exactly the kind of near-identical, hand-copied block described in Section 5.1. The templating syntax introduced below is what turns a call like this into something reusable across every template that needs the same derivation, with only the details that genuinely differ left open as placeholders.

Registered chunks and calls live in an in-memory store attached to the current R session, and `{scriptr}` provides `get_reactable()` to browse the full contents of that store interactively (rendered as a searchable table), which is invaluable when a developer is trying to find out whether a chunk implementing a given piece of logic already exists before writing a new one. Because different packages may want to maintain their own independent chunk libraries, `{scriptr}` supports a **federated store** pattern: `set_scriptr_sources()` lets a downstream package (such as `{admiralrochetools}`, discussed in Section 6) register its own chunk source alongside or instead of the default one, and `get_source_env()` resolves which environment a given lookup should use. The excerpt below, taken from `{admiralrochetools}`’s own template-generation code, shows this in practice: it points `{scriptr}` at its own package namespace *and* at a private, internal environment holding Roche’s confidential chunk and call library, so that the two can be searched together as if they were a single store:

```

set_scriptr_sources(list("admiralrochetools", admiralroche_scriptr_private))

```

5.3 `{scriptr}`’s templating syntax

The templating mechanism is deliberately built on vocabulary borrowed from `{rlang}`’s quasiquotation, since many R developers already have some intuition for it. A chunk or call template can mark a placeholder for single-value substitution with a `++` sigil (mirroring `!!`), and a placeholder for multi-value, “splice in a list” substitution with `+++` (mirroring `!!!`). When a template is instantiated, values are supplied for these placeholders — a single call for a single study parameterisation, or a whole tibble of parameter combinations for row-wise instantiation, which lets a developer generate an entire family of near-identical derivations (for example, one flag derivation per row of a specification table) from a single template definition in a single call.

Two helper functions cover the common cases of building up dynamic text inside a template: `glue_sym()`

constructs a syntactically valid R symbol from string fragments (for example, gluing a variable stem and a suffix into a single identifier), while `glue_char()` builds a dynamic character string, and can lean on `{cli}`'s pluralisation helpers so that, for example, a comment can correctly say “the following variable” versus “the following variables” depending on how many were substituted. Where a literal value needs to be inserted as an R string constant or a vector of string constants rather than interpreted, `quote()` and `quotes()` provide convenience wrappers.

Returning to the AENDTM/AENDT example, `derive_aendt` can be rewritten as a call template with two placeholders: `++prefix`, the variable stem shared by the `*ENDTC` source and the derived `*ENDT`/`*ENDTM` pair, and `++imputation`, the one detail that genuinely differs between the `ADCM` and `ADMH` templates. `glue_sym()` builds the `dtc` argument's symbol from `++prefix` at instantiation time, and `defaults` supplies the value used whenever a caller does not override it:

```
add_call_template(
  "derive_end_dt",
  description = "Derive an end date/time from a *ENDTC variable",
  derive_vars_dt(
    new_vars_prefix = ++prefix,
    dtc = glue_sym("{prefix}ENDTC"),
    highest_imputation = "M",
    date_imputation = ++imputation
  ),
  defaults = list(imputation = "first")
)
```

A single template definition now serves both templates. `ADCM` uses the default “first”-day imputation and need not mention it at all, while `ADMH` overrides it to “last”:

```
insert_calls(
  input = adcm,
  ids = exprs(derive_end_dt(prefix = "AEN")),
  output = adcm
)

insert_calls(
  input = admh,
  ids = exprs(derive_end_dt(prefix = "AEN", imputation = "last")),
  output = admh
)
```

A change to how AENDTM/AENDT should be derived - a new imputation rule, an extra argument, a bug fix - now only ever needs to be made once, in the `derive_end_dt` template, rather than hunted down across every template file that copied it by hand.

5.4 Assembly: from chunks to a finished script

Registering chunks and calls is only half of `{scriptr}`; the other half is turning a selection of them into an actual, well-formatted R script. `insert_chunks()` takes a list of chunk identifiers (conventionally captured as unevaluated expressions via `exprs()`, so that argument names can double as chunk identifiers) and inlines them, in the order given, into the developing script. `insert_calls()` does the equivalent job for calls, and additionally understands the notion of an “input” and an “output” -

letting a developer specify, for instance, that a chain of calls should thread the same dataset through each derivation step in turn, producing a pipe-style sequence rather than a flat list of independent statements. In both cases, ordering is exactly the order supplied by the caller; `{scriptr}` does not attempt to infer dependencies between chunks automatically, which keeps its behaviour predictable and easy to reason about, at the cost of putting the responsibility for correct ordering on whoever is calling it (in practice, either a human developer or, in `{admiralrochetools}`'s case, a purpose-built template generator that already knows the correct domain-specific order).

The top-level entry point that ties this together is `execute()`, which takes an accumulated set of chunk/call insertions together with a target filename and produces a finished `.R` file. Before writing that file, `{scriptr}` passes the assembled code through `{styler}`'s `style_code()`, using a house style configuration, so that the generated script is indistinguishable in formatting from one a Roche programmer would have written by hand — consistent indentation, consistent pipe placement, and no tell-tale artefacts of programmatic generation.

Putting the `load_ae` chunk and the `derive_end_dt` call template from the previous two sections together, generating the relevant section of the ADCM template is a single `execute()` call:

```
execute(
  {
    insert_chunks(ids = exprs(load_ae))

    adcm <- cm
    insert_calls(
      input = adcm,
      ids = exprs(derive_end_dt(prefix = "AEN")),
      output = adcm
    )
  },
  filename = "programs/adam/ad_adcm.R"
)
```

which writes the following, already-formatted script to `ad_adcm.R`, ready for a programmer to open and continue building on:

```
ae <- read_parquet("data/sdtm/ae.parquet")

adcm <- cm
adcm <- admcm %>%
  derive_vars_dt(
    new_vars_prefix = "AEN",
    dtc = AEENDTC,
    highest_imputation = "M",
    date_imputation = "first"
  )
```

The equivalent call for `ADMH` differs only in the arguments passed to `derive_end_dt()`, yet both scripts are generated from the same underlying chunk and call template — precisely the “write it once, generate it everywhere” property that static templates could not offer. Figure 4 shows a further worked example from the talk of a chunk-based template definition alongside the finished, assembled script it produces.

6 Creating a new {admiralrochetools} package to the store template machinery

Once the {scriptr} machinery was in place to help abstract and generalise the {admiralroche} ADaM templates, a decision was made to split {admiralroche} into two:

- {admiralroche} would retain the Roche-specific derivation functions and their documentation
- A new {admiralrochetools} package was created to house all the machinery for the ADaM templates.

The main reason behind this split was of a regulatory nature. Any Roche regulatory submission of R code requires the {admiralroche} package to be submitted too, in order for the ADaM programs to be readable. However, there is no need to submit the ADaM templates, nor the machinery used to create them, since once the code has been generated, the templates are no longer used in the ADaM derivation process. Moving the machinery to {admiralrochetools} meant future submissions would not need to include this package, reducing the size of the submission and the number of packages that need to be documented.

From a user perspective, the only difference was that the `package` argument to `use_ad_template()` needed to be updated to point at {admiralrochetools} instead of {admiralroche}:

```
library(admiral)
library(admiralrochetools)

admiral::use_ad_template(
  "adsl",
  package = "admiralrochetools"
)
```

7 Stage 3 - *Code that writes code that writes code* - Dynamic template generation

7.1 From “assemble what a developer selects” to “ask the study what it needs”

{scriptr} solves the maintenance problem of Chapter 2: a change to shared logic now needs to be made in exactly one chunk, not in twenty templates. But a developer using {scriptr} directly still has to know, for every derivation with more than one valid implementation, which variant this particular study needs, and manually select the right chunks and the right template parameters. {admiralrochetools} closes this final gap by asking the study itself. Every Roche study already has a machine-readable **Data Analysis Plan (DAP) M3 specification** — a structured document (in our case represented as JSON) that records, dataset by dataset and variable by variable, exactly how each ADaM variable should be derived for that study, including the specific sponsor choices (which reference date to use for age, which population definition to use for the intent-to-treat flag, and so on) that Section 3 of this paper described as the genuinely reusable-but-variable middle 10%. {admiralrochetools} reads that specification directly and, wherever it can extract an unambiguous answer from it, uses it; wherever the specification is silent or ambiguous, it asks the programmer a short, targeted question instead.

7.2 Question definitions

Each ADaM domain admiralrochetools supports (at the time of writing: ADSL, ADAE, ADEG, ADEX, ADHY, ADLB, ADRS, ADTTE and ADVS) has an accompanying YAML file of question definitions,

```
- Question: Which subjects do you want to consider for the derivation of `ITTFL`?
Options: `"randomized"`, `"enrolled"`
Default_from_DAP_M3: >
  `"randomized"` if the `ITTFL` derivation contains the text
  `"[DS.DSDECOD] = 'RANDOMIZED'"`, `"enrolled"` otherwise
Notes: ""
```

[illegible]

7.3 From answers to an assembled template

10

example from the talk: the LSTALVDT (last known alive date) derivation, assembled automatically from the programmer’s answers to a handful of short questions.




Figure 4: The assembled LSTALVDT derivation, generated automatically from the answers gathered in Figure 5.

7.4 An audit trail for free

A detail not shown on the slides but worth recording here: every dynamic template-generation session writes an “info file” recording exactly which questions were asked, what the programmer (or the DAP specification) answered, and which package versions produced the result, via `start_input_tracking()/stop_input_tracking()`. This was not a strict requirement when the feature was built, but it has turned out to be a genuinely useful side effect: it gives programmers, reviewers and QC staff downstream a durable record of *why* a generated template looks the way it does, months after the fact, without anyone having had to remember to write that justification down by hand.

7.5 A living tool: version history

{admiralrochetools} has itself evolved quickly since its first release. Version 0.1.0 (“Savannah”) launched with dynamic generation support for five domains — ADSL, ADAE, ADLB, ADRS and ADTTE. Version 0.1.1 (“Ash”) added ADVS support and switched the source of a study’s data-cut date configuration to a newer internal activity-tracking format. Version 0.2.0 (“Bamboo”), the version installed at the time of writing, added ADEX, ADHY and ADEG, bringing the total to eight domain-specific generators, alongside the more general `use_bds_template()` and `use_occds_template()` functions inherited from {admiralroche} in the 1.6.0 split described in Section 5. This cadence — a new domain or two roughly every quarter, tracking the same GDSR release-train codenames used by {admiralroche} itself — gives us reasonable confidence that dynamic generation will, in time, cover the full set of ADaM domains Roche studies routinely need.

8 REFLECTIONS AND FURTHER WORK

Looking back, we do not think we could have jumped straight from static templates (Chapter 1) to fully dynamic, DAP-driven generation (Chapter 3) even if we had known from the outset that this was where we would end up. {scriptr}’s chunk/call abstraction had to exist, be trusted by developers, and prove itself on hand-selected templates first, before it was a credible foundation for a tool that selects and assembles chunks *automatically* on a programmer’s behalf. Each stage of this journey was also, deliberately, a fully **deterministic** piece of automation: a given DAP specification and a given set of question answers will always produce the same generated script, byte for byte. We think

this was the right choice for this problem, precisely because ADaM programming sits inside a heavily regulated, audit-sensitive process where reproducibility and explainability of the generated artefact matter as much as the artefact’s correctness. That said, we have not yet rigorously quantified the maintenance-time or error-rate savings this evolution has produced; doing so properly — for example, by comparing defect rates or review turnaround time on studies built from dynamically generated templates against a matched set built from static ones — remains an open piece of further work, and one we would encourage other companies pursuing a similar path to plan for from the start, rather than treating it as an afterthought as we did.

Looking forward, we see two natural next steps. The first is incremental: extending {admiralrochetools}’s domain coverage and question sophistication, and continuing to fold learnings from real studies back into the chunk library and question definitions as edge cases surface. The second is more speculative, and is where large language models plausibly enter this pipeline for the first time: today, the DAP-specification-parsing helpers described in Section 6 rely on pattern-matching known phrases in free-text derivation descriptions, which is robust but brittle to phrasing the pattern-matcher was not written to expect. An LLM-based reader of the same free-text specification could plausibly resolve a wider range of phrasing with less maintenance of hand-written patterns, while still leaving the final, deterministic assembly step — the part that actually needs to be reproducible and auditable — entirely in {scriptr}’s hands. We would treat this as an enhancement to the specification-reading step only, not a replacement for the deterministic generation engine underneath it.

9 RECOMMENDATIONS

For any organisation considering a similar journey, our main recommendations, in order of how early we think they matter, are:

1. **Do not skip straight to dynamic generation.** Build the reusable-chunk abstraction ({scriptr}’s role in our story) first, on a codebase your developers already trust, before trying to automate the selection of chunks. Trust in the underlying building blocks is what makes the final, more automated layer credible.
2. **Decide submission scope early, and let it drive package boundaries.** Our {admiral-roche}/{admiralrochetools} split (Section 5) happened somewhat after the fact; had we designed for the “does this need to be in a submission’s software environment?” question from the start, we could have made that split earlier and more cleanly.
3. **Make comments and documentation first-class in your assembly engine,** not an afterthought bolted on later. Generated code that a human reviewer cannot understand at a glance will not be trusted, no matter how correct it is.
4. **Keep the generation step deterministic,** and keep an audit trail of what was asked and answered. Both properties cost little to build in from the start and are expensive to retrofit later, and both matter a great deal in a regulated environment.
5. **Read your own specifications before asking your programmers anything.** The single biggest reduction in manual effort in our pipeline did not come from generating code faster — it came from not asking a question the DAP specification had already answered.

10 CONCLUSION

The journey from static ADaM templates, through {scriptr}’s reusable chunk abstraction, to {admiral-rochetools}’s fully dynamic, specification-driven template generation reflects a fairly ordinary software engineering instinct — notice duplication, abstract it, then automate the selection of the abstraction —

applied to a domain, regulated clinical trial programming, where that instinct is not always the first one people reach for. We hope that by setting out both the narrative (in the accompanying talk) and the internal mechanics (in this paper) we have given other sponsor companies a genuine starting point for making the same journey, whether or not they choose the same tools to make it.

11 ACKNOWLEDGEMENTS

We are especially grateful to Stefan Bundfuss, the lead developer of both `{admiralroche}` and `{scriptr}`, whose design work is the foundation this entire talk and paper describe. We also thank our wider Roche statistical programming community for continuing to build, test and refine the templates and chunk libraries described here.

12 RECOMMENDED READING

- Farrugia, R. (2024). *{admiral}* - from open source through to company implementation. PHUSE EU Connect 2024, Paper OS05.
- Mancini, E. (2024). *Two hats, one noggin: Working as a developer and as a user of the {admiral} R package for creating ADaMs*. PHUSE EU Connect 2024, Paper OS07.
- The pharmaverse family of R packages: <https://pharmaverse.org>
- `{admiral}` package documentation and Programming Strategy vignette: <https://pharmaverse.github.io/admiral/>

13 CONTACT INFORMATION

Your comments and questions are valued and encouraged. Contact the authors at:

Gordon Miller Roche Products Ltd Welwyn Garden City, UK Email: gordon.miller@roche.com

Edoardo Mancini Roche Products Ltd Welwyn Garden City, UK Email: edoardo.mancini@roche.com

Any trademarks mentioned in this paper are the property of their respective owners.